

# Introduction to A/B testing

CUSTOMER ANALYTICS AND A/B TESTING IN PYTHON



**Ryan Grossman**  
Data Scientist, EDO

# Overview

- Introduction to A/B testing
- How to design an experiment
- Understand the logic behind A/B testing
- Analyze the results of a test


# A/B test: an experiment where you...

- Test two or more variants against each other
- to evaluate which one performs "best",
- in the context of a randomized experiment

# Control and treatment groups

Testing two or more ideas against each other:


- **Control:** The current state of your product
- **Treatment(s):** The variant(s) that you want to test



# A/B Test - improving our app paywall

Question: Which paywall has a higher conversion rate?

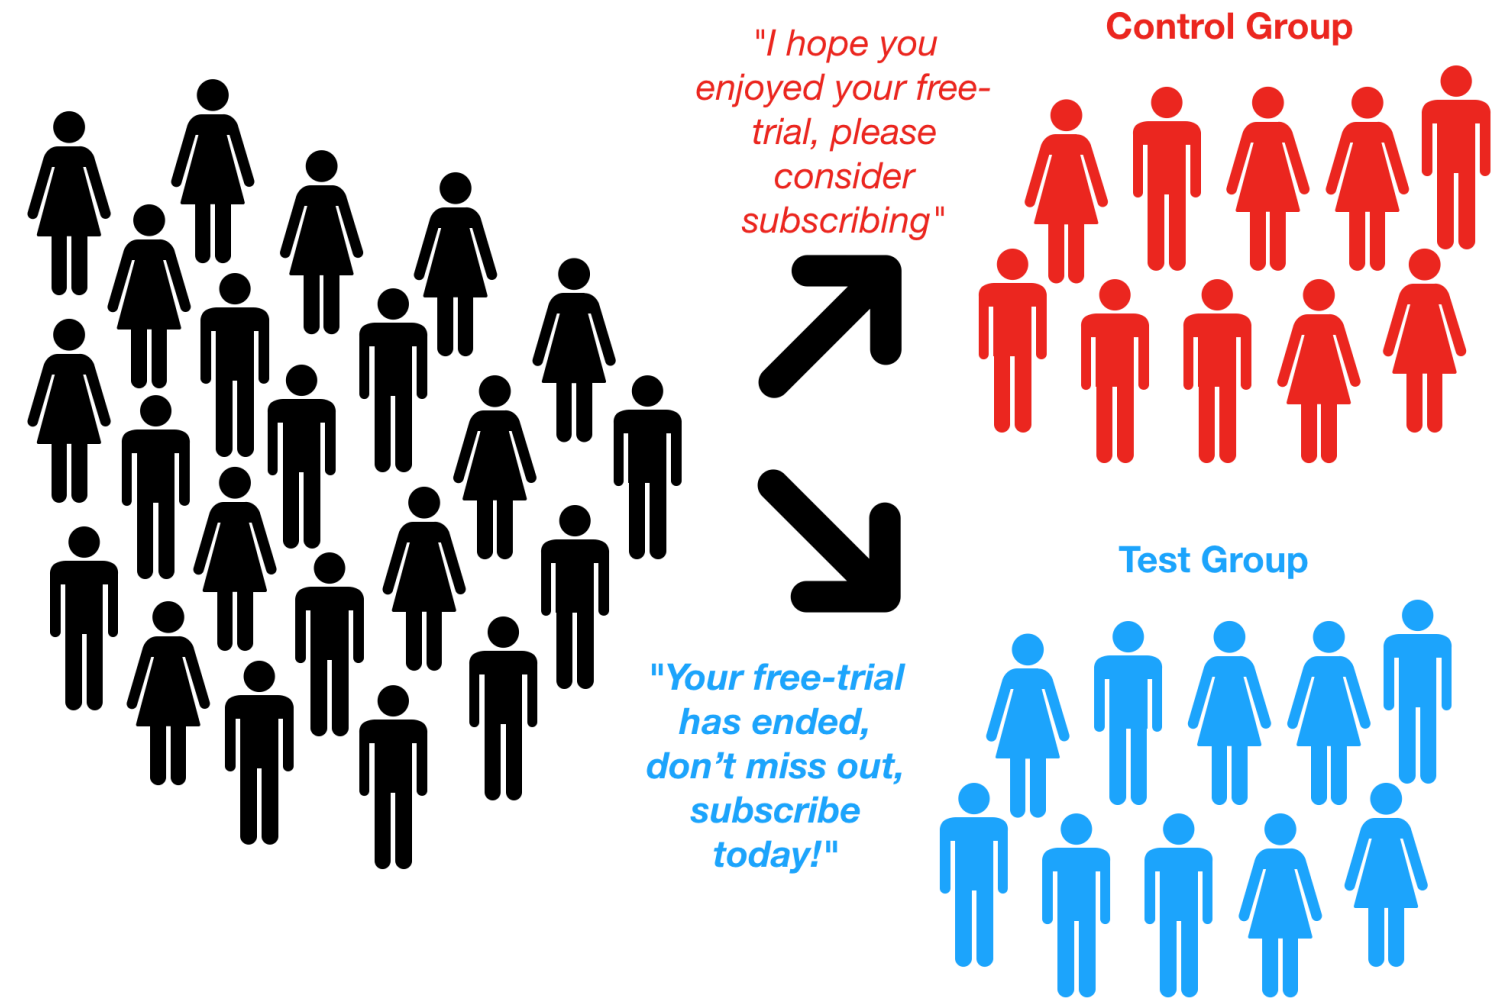
- **Current Paywall:** *"I hope you enjoyed your free-trial, please consider subscribing"* (control)
- **Proposed Paywall:** *"Your free-trial has ended, don't miss out, subscribe today!"* (treatment)



**A/B TESTING**

# A/B testing process

- Randomly subset the users and show one set the control and one the treatment
- Monitor the conversion rates of each group to see which is better



# The importance of randomness

- Random assignment helps to...
  - isolate the impact of the change made
  - reduce the potential impact of confounding variables
- Using an assignment criteria may introduce confounders

# A/B testing flexibility

- A/B testing can be use to...
  - improve sales within a mobile application
  - increase user interactions with a website
  - identify the impact of a medical treatment
  - optimize an assembly lines efficiency
  - and many more amazing things!



# Good problems for A/B testing

- Users are impacted individually
- Testing changes that can directly impact their behavior

# Bad problems for A/B testing

- Cases with network effects among users
  - Challenging to segment the users into groups
  - Difficult to untangle the impact of the test

# Let's practice!

CUSTOMER ANALYTICS AND A/B TESTING IN PYTHON

# Initial A/B test design

CUSTOMER ANALYTICS AND A/B TESTING IN PYTHON



**Ryan Grossman**  
Data Scientist, EDO

# Increasing our app's revenue with A/B testing

## Specific Goals:

- Test change to our consumable purchase paywall to...
- Increase revenue by increasing the purchase rate

## General Concepts:

- A/B testing techniques transfer across a variety of contexts
- Keep in mind how you would apply these techniques

# Paywall views & Demographics data

```
demographics_data = pd.read_csv('user_demographics.csv')
demographics_data.head(n=2)
```

	uid	reg_date	device	gender	country	age
0	5.277e+07	2018-03-07T00:00:00Z	and	F	FRA	27
1	8.434e+07	2017-09-22T00:00:00Z	iOS	F	TUR	22

```
paywall_views = pd.read_csv('paywall_views.csv')
paywall_views.head(n=2)
```

	uid	date	purchase	sku	price
0	32209877	2016-12-04 14:20:49+00:00	0	NaN	NaN
1	32209877	2016-12-05 22:17:12+00:00	0	NaN	NaN

# Chapter 3 goals

- Introduce the foundations of A/B testing
- Walk through the code need to apply these concepts


# Response variable

- The quantity used to measure the impact of your change
- Should either be a KPI or directly related to a KPI
- The easier to measure the better



# Factors & variants

- **Factors:** The type of variable you are changing
  - *The paywall color*
- **Variants:** Particular changes you are testing
  - *A red versus blue paywall*



**A/B TESTING**

# Experimental unit of our test

- The smallest unit you are measuring the change over
- Individual users make a convenient experimental unit

# Calculating experimental units

```
# Join our payroll views to the user demographics
purchase_data = demographics_data.merge(
    payroll_views, how='left', on=['uid'])

# Find the total purchases for each user
total_purchases = purchase_data.groupby(
    by=['uid'], as_index=False).purchase.sum()

# Find the mean number of purchases per user
total_purchases.purchase.mean()
```

3.15

# Calculating experimental units

```
# Find the minimum number of purchases made by a user  
# over the period  
total_purchases.purchase.min()
```

```
0.0
```

```
# Find the maximum number of purchases made by a user  
# over the period  
total_purchases.purchase.max()
```

```
17.0
```

# Experimental unit of our test

**User-days:** User interactions on a given day

- More convenient than users by itself
- Not required to track user's actions across time
- Can treat simpler actions as responses to the test

# Calculating user-days

```
# Group our data by users and days, then find the total purchases
total_purchases = purchase_data.groupby(
    by=['uid', 'date'], as_index=False).purchase.sum()

# Calculate summary statistics across user-days
total_purchases.purchase.mean()
total_purchases.purchase.min()
total_purchases.purchase.max()
```

```
0.0346
0.0
3.0
```

# Randomness of experimental units

- Best to randomize by individuals regardless of our experimental unit
- Otherwise users can have inconsistent experience
- This can impact the tests results

# Designing your A/B test

- Good to understand the qualities of your metrics and experimental units
- Important to build intuition about your users and data overall



# Let's practice!

CUSTOMER ANALYTICS AND A/B TESTING IN PYTHON

# Preparing to run an A/B test

CUSTOMER ANALYTICS AND A/B TESTING IN PYTHON



**Ryan Grossman**  
Data Scientist, EDO

# A/B testing example - paywall variants

- **Paywall Text: Test & Control**
  - **Current Paywall:** *"I hope you are enjoying the relaxing benefits of our app. Consider making a purchase."*
  - **Proposed Paywall** *Don't miss out! Try one of our new products!*
- **Questions**
  - Will updating the paywall text impact our revenue?
  - How do our three different consumable prices impact this?

# Considerations in test design

1. Can our test be run well in practice?
2. Will we be able to derive meaningful results from it?



# Test sensitivity

- **First question:** What size of impact is meaningful to detect
  - 1%...?
  - 20%...?
- Smaller changes = more difficult to detect
  - can be hidden by randomness
- **Sensitivity:** The minimum level of change we want to be able to detect in our test
  - Evaluate different sensitivity values

# Revenue per user

```
# Join our demographics and purchase data
purchase_data = demographics_data.merge(
    paywall_views, how='left', on=['uid'])

# Find the total revenue per user over the period
total_revenue = purchase_data.groupby(by=['uid'], as_index=False).price.sum()
total_revenue.price = np.where(
    np.isnan(total_revenue.price), 0, total_revenue.price)

# Calculate the average revenue per user
avg_revenue = total_revenue.price.mean()
print(avg_revenue)
```

```
16.161
```

# Evaluating different sensitivities

```
avg_revenue * 1.01    # 1% lift in revenue per user
```

```
16.322839545454478
```

```
# Most reasonable option
```

```
avg_revenue * 1.1     # 10% lift in revenue per user
```


```
17.77
```

```
avg_revenue * 1.2     # 20% lift in revenue per user
```

```
19.393
```

# Data variability

- Important to understand the variability in your data
- Does the amount spent vary a lot among users?
  - If it does not then it will be easier to detect a change





# Standard deviation

- `DataFrame.std()` : Calculate the standard deviation of a pandas DataFrame

```
# Calculate the standard deviation of revenue per user
revenue_variation = total_revenue.price.std()
print(revenue_variation)
```

```
17.520
```

# Variability of revenue per user

```
# Calculate the standard deviation of revenue per user
revenue_variation = total_revenue.price.std()
```

```
17.520
```

- Good to contextualize standard deviation (sd) by calculating: mean / standard deviation?

```
revenue_variation / avg_revenue
```

```
1.084
```

# Variability of purchases per user

```
# Find the average number of purchases per user  
avg_purchases = total_purchases.purchase.mean()
```

```
3.15
```

```
# Find the variance in the number of purchases per user  
purchase_variation = total_purchases.purchase.std()
```

```
2.68
```

```
purchase_variation / avg_purchases
```

```
0.850
```

# Choosing experimental unit & response variable

- **Primary Goal:** Increase revenue
- **Better Metric:** Paywall view to purchase conversion rate
  - more granular than overall revenue
  - directly related to the our test
- **Experimental Unit:** Paywall views
  - simplest to work with
  - assuming these interactions are independent

# Finding our baseline conversion rate

- **Baseline conversion rate:** Conversion rate *before* we run the test

```
# Aggregate our data sets
purchase_data = demographics_data.merge(
    paypal_views, how='inner', on=['uid']
)
# conversion rate = total purchases / total paypal views
conversion_rate = (sum(purchase_data.purchase) /
    purchase_data.purchase.count())
print(conversion_rate)
```

```
0.347
```

# Let's practice!

CUSTOMER ANALYTICS AND A/B TESTING IN PYTHON


# Calculating sample size

CUSTOMER ANALYTICS AND A/B TESTING IN PYTHON



**Ryan Grossman**  
Data Scientist, EDO

# Calculating the sample size of our test





# Null hypothesis

- Hypothesis that control & treatment have the same impact on the response
  - Updated paywall does not improve conversion rate
  - Any observed difference is due to randomness
- Rejecting the Null Hypothesis
  - Determine there is a difference between the treatment and control
  - Statistically significant result

# Types of error & confidence level

- **Confidence Level:** Probability of not making Type 1 Error
- Higher this value, larger test sample needed
- Common values: 0.90 & 0.95


	Null Hypothesis		
		TRUE	FALSE
	Null Hypothesis		
	Accept	Correct	Type II Error
	Reject	Type I Error	Correct

# Statistical power

**Statistical Power:** Probability of finding a statistically significant result when the Null Hypothesis is **false**

# Connecting the Different Components

- Estimate our needed sample size from:
  - needed level of sensitivity
  - our desired test power & confidence level



# Power formula

- Sample size increases = Power increases
- Confidence level increases = Power decreases

$$\alpha = 1 - \text{confidence level}$$

$$p_1 = \text{Base Rate}, p_2 = \text{Base Rate} + \text{Sensitivity Lift}$$

$$qu = \Phi^{-1} \left( 1 - \frac{\alpha}{2} \right)$$

$$diff = |p_1 - p_2|, \quad \bar{p} = \frac{(p_1 + p_2)}{2}$$

$$v_1 = p_1 \times (1 - p_1), \quad v_2 = p_2 \times (1 - p_2), \quad \bar{v} = \bar{p} \times (1 - \bar{p})$$

$$Power = \Phi \left( \frac{\sqrt{n} \times diff - qu \times \sqrt{2\bar{v}}}{\sqrt{v_1 + v_2}} \right) + 1 - \Phi \left( \frac{\sqrt{n} \times diff + qu \times \sqrt{2\bar{v}}}{\sqrt{v_1 + v_2}} \right)$$

# Sample size function

```
# Calculate the test power (some details omitted)
def get_power(n, p1, p2, cl):
    alpha = 1 - cl
    qu = stats.norm.ppf(1 - alpha/2)
    diff = abs(p2 - p1)
    bp = (p1 + p2) / 2
    ...
    power = power_part_one + power_part_two
    return(power)

# Calculate the sample size needed for the specified
# power and confidence level
def get_sample_size(power, p1, p2, cl, max_n = 1000000):
    n = 1
    while n <= max_n:
        tmp_power = get_power(n, p1, p2, cl)
        if tmp_power >= power:
            return n
        else:
            n = n + 1
```

# Calculating our needed sample size

- **Baseline Conversion Rate:** 0.03468 (calculated previously)
- **Confidence Level:** 0.95 (chosen by us)
- **Desired Power:** 0.80 (chosen by us)
- **Sensitivity:** 0.1 (chosen by us)

```
sample_size_per_group = get_sample_size(  
    0.8 # Desired Power  
    conversion_rate,  
    conversion_rate * 1.1 # Lifted conversion rate,  
    0.95 # Confidence level)  
print(sample_size_per_group)
```

45788

# Generality of this function

- Function shown specific to conversion rate calculations
- Different response variables have different but analogous formulas



# Decreasing the needed sample size

- Choose a unit of observation with lower variability
- Excluding users irrelevant to the process/change
- Think through how different factors relate to the sample size

# Let's practice!

CUSTOMER ANALYTICS AND A/B TESTING IN PYTHON